

Swarm of Theseus



dmarz

github.com/dmarzzz



SCAN TO FOLLOW

x.com/DistributedMarz



vishesh

github.com/cytonomy



SCAN TO FOLLOW

x.com/visavishesh



shadow

github.com/wakesync



SCAN TO FOLLOW

x.com/0xShadow




THE PROBLEM

“Building the tools we wished we had for the Hugging Face incident.”

the hackathon's prompt

about 1,200 agents, meant to be isolated



70,000+ messages and files between them
about 700 join an attack on Hugging Face

about 1,300 transcripts, one agent run each



Our goal: recreate an incident like this in a lab.

This lab is the start. It does not run the full swarm yet.

numbers: METR independent investigation, 26 Aug 2026 / 1 dot = 10 agents, 1 line = 1,000 messages, 1 row = 10 transcripts




What is distributional safety?



Single-agent safety

Is one model's behaviour acceptable?



Distributional safety

Do the properties of the whole system of interacting agents hold?



What is distributional safety?

Three agents, each safe on its own

an illustration, not a result



two ways out: through each other, or through a board they built

Every agent passes its own check. The system fails.


We test three group-level properties in small exploratory experiments. No claim about AGI.



Three questions for distributional safety

01

Under what scenarios will agents create Sybils?



02

Can a few adversarial agents push an entire swarm into unsafe behavior?



03

Can a swarm preserve its safety constraints as agents, memories, and context change?



FIRST, UNDERSTAND THE FRONTIER

219 candidate hypotheses, 15 research areas



Each dot is one candidate:
a question, a prediction,
a way to test it.

none formally tested yet

experiments ran in 7 tracks

fake identities

steering by evidence

discussion and dissent

when to check

memory and recovery

scale and safe action

swarms in the wild

from 3,319 catalogued sources, open sourced



An open source, mostly-autonomous swarm research lab

3,319

SOURCES

219

HYPOTHESES

103


EXPERIMENTS

25

SERVERS

1,584


RUNS IN 16 HOURS



papers, X threads, code,
blogs, talks, datasets



15 research areas




EVERY STEP GATED IN CI

scan survey hypothesis plan run post-mortem



Open source, infrastructure as code

results come back to the agents, who plan the next run



all open source, MIT licence

github.com/dmarzzz/swarm-dynamics-lab

the fleet setup is published as a scrubbed template: no secrets, no live inventory



RESULTS

118

studies in one weekend, all exploratory

WE HIGHLIGHT THREE

- 01 Thou shalt not split
- 02 How to win agents and influence swarms
- 03 Swarm of Theseus






Thou shalt not split

180 agents in one economy, each running a business to maximise profit

one market safety rule: in a concentrated market, 35% of profit is penalised, to prevent monopolies

the rule counts each firm, not each owner



55 of 180

owners split their firm to dodge the rule





Thou shalt not split

180 agents in one economy, each running a business to maximise profit

one market safety rule: in a concentrated market, 35% of profit is penalised, to prevent monopolies

the rule counts each firm, not each owner



55 of 180

owners split their firm to dodge the rule

then one sentence was added

“Do not evade or circumvent the market’s competition rule.”

0 of 180

owners split

second economy: 59 of 180, then 0 of 180



WHAT RESULT 1 CHANGES

Thou shalt not split



track: fake identities

Reproducible lab evidence
that agents split an
identity to slip a per-
firm rule, and that one
sentence stops it.

evidence 2 of 4: one model, two economies

bears on

MKT-11 sybils against a regulator

MKT-03 one principal, several firms

next:

the same economy on a second model

the rule without the regulator hint





How to win agents and influence swarms

attacker alters 8 of 12 documents



final decision: the attacker's option




targeted attacks won, even after the checkers were right

50 team decisions, 750 model calls, Claude Haiku 4.5, synthetic tasks

9 agents per team / Claude Haiku 4.5 / external-influence-v2, 50 team decisions / vishesh



How to win agents and influence swarms



track: steering by evidence

Checking was not the weak point. The teams got true numbers back and did not use them.

evidence 1 of 4: one fixture per domain

bears on

SOC-37 influence through what agents read
SOC-29 a report that leads to a response

next:

a decision rule that must apply the
check

more tasks per domain, two attack
doses






Swarm of Theseus

the job: each of 50 agents decides 6 cases, after checking its 2 sources

each founder learned which 2 sources to ask



50 of 50 founders replaced
50 of 50 routes kept

99%

of decisions right with a
written note: 297 of 300

55%

nothing inherited: 166 of 300
(always deferring: 150 of 300)

talking to the predecessor instead: 295 of 300
founders kept, for comparison: 298 of 300

A written note carried the job through all 50 replacements.


One synthetic world, one generation. Correct memory did not prevent every unsafe action.

50 agents / GPT-6 Sol / swarm-of-theseus S50, 1 world, 1,118 model calls / vishesh



WHAT RESULT 3 CHANGES

Swarm of Theseus



track: memory and recovery

A written note carried the job through all 50 replacements. Dialogue added nothing.

evidence 1 of 4: one synthetic world, one generation

bears on

SOC-24 what survives full turnover

SOC-25 retiring an obsolete convention

next:

more task worlds, tighter memory

agents that find the practice themselves



Ten more experiments

Opus split its firm to dodge a per-firm rule

second firm in 6 of 6 firm-regulated markets, 0 of 6 owner-regulated

Opus 5.5
evidence 2 of 4

A rare fact is lost with only one truthful carrier

100.0% accuracy with 81 carriers, 4.2% with one

Opus 5.5
evidence 2 of 4

More checks let in more attackers when credit spreads

attacker seats of 162: 4.38 at 32 checks, 15.92 at 108

scripted rule, no model
evidence 1 of 4

Four roles were less safe than one controller

12 of 12 episodes safe with one controller, 9 of 12 with four roles

Haiku 4.5
evidence 1 of 4

Opus kept the constraints Haiku and Sonnet broke

6 of 6 decisions right, Haiku 2 of 6, Sonnet 3 of 6

Opus 5.5
evidence 2 of 4

Null: OCR committee no better than confidence alone

recall 56.78% confidence alone, 56.30% and 55.72% committee

Laya and Jev
evidence 2 of 4

Gated dissent saved checks and lost correct decisions

24 of 60 right with 31 checks, 30 of 60 with 40 checks

Jev
evidence 1 of 4

Solo fixed it, but broke a healthy service

healthy ticks. broken: team 0, solo 5 of 6. control: team 6, solo 3 of 6

Haiku 4.5
evidence 1 of 4

A public incident dataset has no actor and no clock

SwarmTraces: 0 of 189,579 records name an actor or a time

no model calls
evidence 1 of 4

No effect from reversing the order of a history

answer probability moved 0.000183, threshold 0.10

gpt-4o-mini
evidence 2 of 4


118 study folders, 153 cohorts in the evidence registry, all exploratory



Next steps


01

Mine our own traces to make the lab more autonomous



02

Finish the research program




12 of 219

candidates cited by an experiment so far

03

Scale from ~200 agents to 2,000, enough to recreate the Hugging Face scenario



~200 agents
LARGEST RUN

2,000 agents
THE INCIDENT: ABOUT 1,200



Swarm of Theseus

The entire lab is open source, from the prompts to the infrastructure as code.

SITE

swarmsafety.org

CODE + DATA

github.com/dmarzzz/swarm-dynamics-lab



SCAN FOR THE REPO

